

# JUECHU DONG

✉ joydong@umich.edu

⌚ joyddd.github.io

## SUMMARY

Juechu (Joy) Dong is a Ph.D. candidate at the University of Michigan, advised by Prof. Satish Narayanasamy. She studies emerging technologies in computer architecture and systems, with a focus on confidential computing and GPU kernel optimizations. Her research seeks to democratize kernel customization by building flexible and adaptive infrastructure for mapping novel algorithms to GPU hardware.

## EDUCATION

<b>University of Michigan - Ann Arbor</b> <i>Ph.D., Computer Science and Engineering</i> Advisor: Prof. Satish Narayanasamy	(exp.) 2027
<b>Shanghai Jiao Tong University</b> <i>B.S., Computer Engineering</i>	2022
<b>University of Michigan - Ann Arbor</b> <i>B.S.E., Computer Engineering, Summa Cum Laude</i> GPA: 3.99/4.00	2022

## SELECTED HONORS

<b>MLCommons ML and Systems Rising Star</b> <i>Selected as one of 38 junior researchers worldwide fostering potential in ML and Systems research.</i>	2025
<b>Meta 2024 Internship Project Spotlight: FlexDecoding</b> <i>Awarded as one of 3 outstanding internship projects each year</i>	2024
<b>Rackham Doctoral Intern Fellowship</b>	2025
<b>Rackham International Student Fellowship (12,990 USD)</b>	2023-24

## INDUSTRY EXPERIENCE

<b>PyTorch group, Meta Inc.</b> <i>Research Scientist Intern</i> - Contribute to TorchInductor & Helion DSL. - Develop new techniques in PyTorch compiler with a focus on GPU performance optimization. - Design GPU programming language for fast, flexible, and easy-to-use ML kernel authoring. - Research new techniques for high-performance distributed GPU communication. - Engage in the open source community to identify user needs and promote new features.	2024,25
<b>NVIDIA</b> <i>Deep Learning Architect Intern</i> - Modeled and analyzed next-gen GPU memory features, including distributed shared memory, asynchronous transaction barrier, and Tensor Memory Accelerator (TMA). - Designed and improved compiler abstractions and programming interfaces for GPU domain-specific languages. - Developed modular framework for CuTe DSL decoupling algorithmic logic from GPU scheduling, enabling flexible kernel composition and ML framework integration.	2022,25

## PUBLICATIONS

- [1] Juechu Dong\*, Boyuan Feng\*, Driss Guessous\*, Yanbo Liang\*, Horace He. "Flex Attention: A Programming Model for Generating Optimized Attention Kernels". In *Proceedings of Machine Learning and Systems 7. (MLSys '25)* 2025.
- Develop a novel compiler-driven programming model that allows implementing the majority of attention variants in a few lines of idiomatic PyTorch code.

- Optimize customizable attention kernels to provides 1.1x - 1.3x speedup compared to FlashAttn2 by lowering customizable attention into a fast Triton kernel + taking advantage of sparsity.
  - Adapt FlexAttention to efficiently support decoding, GQA and PagedAttention.
- [2] Juechu Dong, Jonathon Rosenblum, Satish Narayanasamy. "Toleo: Scaling Freshness to Tera-scale Memory Using CXL and PIM". In *Proceedings of the 29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 4. (ASPLOS '24)* 2024.
- Scale trusted memory size from hundreds of MB to tens of TB by expanding the span of trusted from a single trusted processor to an entire platform including intelligent memories.
  - Design a new scheme of freshness protection that reduces the space requirement by 50x.
  - Reduce deployment cost by spacing sharing one intelligent memory device among multiple CPUs.
- [3] Juechu Dong\*, Xueshen Liu\*, Harisankar Sadasivan, Sriranjani Sitaraman, Satish Narayanasamy. "mm2-gb: GPU Accelerated Minimap2 for Long Read DNA Mapping". In *Proceedings of the 15th ACM International Conference on Bioinformatics, Computational Biology, and Health Informatics. (BCB '24)*<sup>1</sup> 2024.
- Accelerate computational intensive chaining step in the state-of-art long sequence mapping tool minimap2 using AMD GPU by 2.57-5.33x.
  - Optimize towards ultra long reads of 100k+ to accommodate genome sequencing technology trend.
  - Develop adaptive GPU scheduling algorithm to balance highly heterogeneous workload.
- [4] Jonathon Rosenblum, Juechu Dong, Satish Narayanasamy. "SECRET-GWAS: Confidential Computing for Population-Scale GWAS". In *Nature Computer Science*. 2025.
- Develop a thousand-core platform on Azure Confidential Computing to conduct multi-institutional GWAS on millions of patients in less than a minute.
  - Adapt Spark-based Hail genomic analysis framework to run on TEE under obliviousness requirement.
  - Parallelize GWAS computation on 1k cores to achieve near linear speedup.

## PROJECTS (Work in Progress)

---

**Helion: Python-embedded Domain-Specific Language (DSL) for High-Performance ML Kernels** 2025 – Present

- Design and extend higher-level DSL for efficient distributed ML kernels authoring.
- Investigate into in-kernel communication and computation fusion for GPUs (PyTorch conf poster: Kraken)
- Implement cross-device communication/sync abstraction and lowering path in Helion.

**GCnC: GPU Kernel DSL Abstraction Briding Framework** 2025 - present

- Designed abstraction framework bridging domain expertise and low-level GPU optimization in CuTe DSL.
- Developed compiler transformation preserving scheduling semantics for high-level kernel customization.
- Implement abstraction lifting data flow graph generator for CuTe DSL, enabling visualization & sync sanitization.
- Targetting submission to ASPLOS 2027

## SKILLS

---

**Programming Languages:** C/C++, CUDA, python, Triton, verilog

**Technologies/Frameworks:**

*ML Framework:* PyTorch (TorchInductor, TorchDistributed)

*GPU DSL:* Helion, CuTeDSL, cuTile, Triton, ThunderKittens, CUDA, PTX

*GPU Tuning:* nsight-compute/nsight-sys, omniperf/omnitrace/rocprof

*Simulation Software:* SniperSim, DRAMSim, pinplay

*Confidential Computing:* Open Enclave SDK, Intel SGX

**Architectures:** AMD CDNA2 Instinct GPU, NVIDIA Hopper/Blackwell, Intel Xeon Phi, Out-of-order CPU

**Compiler Design:** MLIR, CUTLASS IR, TileIR, Helion IR

---

<sup>1</sup>ACM-BCB is the flagship conference of the ACM SIGBio.